BabelDOC Smoke Test

This document exists to verify that the BabelDOC pipeline can parse a PDF, analyse its layout, and write translated output.
The transformer architecture replaced recurrence with attention, which made training on long sequences dramatically cheaper and parallelisable.
Every paragraph in this file is ordinary body text, so the layout analyser should treat each block as a paragraph and translate it as a whole.

Results

We evaluated the model on three benchmarks and report the mean of five runs. Performance improved steadily with model size, though the gains flattened beyond roughly ten billion parameters.

The ablation study confirms that both components contribute to accuracy.